



UNIVERSIDADE ESTADUAL DE CAMPINAS
Faculdade de Engenharia Civil, Arquitetura e Urbanismo

IC639: Métodos Numéricos para Engenharia Civil

List 4

Iterative Methods for Linear Systems of Equations

Student:

Carlos Henrique Chama Puga - 195416

Advisors:

Porf. Dr. Philippe Devloo

Dr. Giovane Avancini

Campinas

2024

Contents

1	Introduction	3
2	Modification in the SparseMatrix Class	3
2.1	Parsing a Sparse Matrix from a File	3
2.2	Getting the Diagonal Elements	5
2.3	Performing the Inner Product between the Lower and Upper Tri- angles and a Vector	5
3	The IndirectSolver Class	7
3.1	Solve Method	8
3.2	Conjugate Gradient Method	9
3.3	The Jacobi Preconditioner	10
3.4	The SSOR Preconditioner	11
4	Comparison of the Conjugate Gradient Method	12
5	Conclusions	15
	References	15
A	GitHub Repository	15

1 Introduction

2 Modification in the SparseMatrix Class

Before we properly dive into the implementation of Iterative Methods for solving linear systems, small modifications in the previously implemented SparseMatrix class are required. SparseMatrix class was implemented in List 3 (see Appendix A) and the following modifications are necessary:

- A method to parse a sparse matrix from a file;
- A method to get the diagonal elements;
- A method to perform the inner product between the vector made by the elements of the lower triangle by a vector of the same size;
- The same method as above, but for the upper triangle;

2.1 Parsing a Sparse Matrix from a File

The ParseFromFile method is responsible for reading a file and parsing its content into a SparseMatrix object. Code 1 shows the implementation of this method.

```
1 def ParseFromFile(self, file:str)->None:
2     row = 0
3     col = 0
4     end_row = False
5     self.ptr.append(0)
6
7     with open(file, "r") as f:
8         lines = f.readlines()
9
10        for line in lines:
11            line = line.strip("=" { , \n")
```

```
12
13
14     for element in line.split(","):
15         if "}" in element:
16             element = element.strip("}")
17             end_row = True
18
19         try: element = float(element)
20         except ValueError: continue
21
22         if element:
23             self.data.append(element)
24             self.cols.append(col)
25
26         col += 1
27
28         if end_row:
29             self.ptr.append(len(self.data))
30             row += 1
31             col = 0
32             end_row = False
33
34     self.size = row
```

Code Listing 1: ParseFromFile method implementation

This method sets the cols, data and ptr vector required to represent a sparse matrix. In line 15, the code checks if the element has a closing bracket, which means that the row has ended. If so, the end_row flag is set to True and the row number is incremented.

Lines 19 and 20 try to convert the element (at this point a string) to a float. If the conversion is not successful, the code continues the iteration. If the conversion is successful, the element is checked to verify whether it is null. If the element is non-null, it is appended to the data vector and its column to the cols vector.

Lines 28 to 32 update the row and reset the column counter in case the row has ended. The size of the matrix is set to the number of rows found in the file.

2.2 Getting the Diagonal Elements

During the Jacobi method, it is necessary to retrieve the matrix's diagonal elements to perform the update of the residual and the solution. The `GetDiagonal` method is implemented to perform this task. Code 2 shows the implementation of this method.

```
1 def GetDiagonal(self)->np.array:
2     diag = np.zeros(self.size)
3     for i in range(self.size):
4         diag[i] = self.FindAij(i, i)
5
6     return diag
```

Code Listing 2: `GetDiagonal` method implementation

A vector called `diag` is returned by the method. This vector is filled with the diagonal elements, which are obtained by calling the `FindAij` method (see List 3 implementation in Appendix A) for each element in the diagonal.

2.3 Performing the Inner Product between the Lower and Upper Triangles and a Vector

Finally, the last modifications are the `InnerProductLowerRows` and `InnerProductUpperRows` methods. These methods are responsible for performing the inner product between the lower and upper triangles of the matrix and a vector, a procedure required in the Gauss-Seidel method. Code 3 shows the implementation of these methods.

```
1 def InnerProductLowerRows(self, vector:np.array, row:int)->np.array:
2     nelem = self.ptr[row+1] - self.ptr[row]
3     ntotal = self.ptr[row]
```

```
4
5     sum = 0
6     for j in range(ntotal, ntotal + nelem):
7         if self.cols[j] >= row: continue
8
9         sum += self.data[j] * vector[self.cols[j]]
10
11     return sum
12
13 def InnerProductUpperRows(self, vector:np.array, row:int)->np.array:
14     nelem = self.ptr[row+1] - self.ptr[row]
15     ntotal = self.ptr[row]
16
17     sum = 0
18     for j in range(ntotal, ntotal + nelem):
19         if self.cols[j] <= row: continue
20
21         sum += self.data[j] * vector[self.cols[j]]
22
23     return sum
```

Code Listing 3: InnerProductLowerRows and InnerProductUpperRows methods implementation

The only difference between both methods is that, while the first one searches for elements in the lower triangle (if `self.cols[j] >= row` - Line 7), the second one searches for elements in the upper triangle (if `self.cols[j] <= row` - Line 15).

Since the Gauss-Seidel method can be performed backward and forward, both implementations are required. Another method that employs both implementations is the Symmetric Successive Over-Relaxation (SSOR) method, which performs both forward and backward Gauss-Seidel iterations.

With these modifications done, we can now implement the Iterative Methods for solving linear systems.

3 The IndirectSolver Class

This list's main goal is to implement and verify the performance of Iterative Methods for solving linear systems. Herein are implemented and compared the Conjugate Gradient (CG) and the Preconditioned Conjugate Gradient (PCG) methods and as preconditioners, the Jacobi and the SSOR methods. For simplicity, the preconditioned conjugate gradient with Jacobi and with SSOR are here referred to as CG-J and CG-SSOR, respectively.

The IndirectSolver class is responsible for implementing these methods. The class constructor is shown in Code 4.

```
1 @dataclass
2 class IndirectSolver:
3     A: SparseMatrix
4     rhs: np.ndarray
5     niter: int
6     omega: float = 1.0
7     method: callable = None
8
9     resnorm: list[float] = field(init=False, default_factory=list)
10
11     p_k: np.array = field(init=False, default=list)
12     res_k: np.array = field(init=False, default=list)
13
14     preconditioner: callable = field(init=False, default=None)
15     z: np.array = field(init=False, default=None)
16     z_k: np.array = field(init=False, default=None)
```

Code Listing 4: IndirectSolver class constructor

As parameters, the class receives the sparse matrix A, the right-hand side vector rhs and the number of iterations niter. The relaxation parameter omega and the method to be used are optional parameters, although the method is set after. The resnorm list stores the norm of the residual at each iteration. The p_k,

res_k , z , and z_k vectors are used to store the search direction, the residual, and the preconditioned vectors. A preconditioner might be set to the solver.

Hereafter, the methods implemented in the class are presented.

3.1 Solve Method

The solve method is a general function that calls the method set in the constructor to solve the linear system. Code 5 shows the implementation of this method.

```

1 def Solve(self) -> None:
2     if not self.method:
3         raise ValueError("Method not set")
4
5     sol = ZEROS(len(self.rhs))
6     res = self.rhs - self.A.Multiply(sol)
7
8     if self.method == self.ConjugateGradient:
9         if not self.preconditioner:
10             self.p_k = res.copy()
11         else:
12             self.z, _ = self.preconditioner(ZEROS(self.A.size), res)
13             self.p_k = self.z.copy()
14             self.z_k = self.z.copy()
15
16     self.res_k = res.copy()
17
18     self.resnorm = [NORM(res)]
19     for i in range(self.niter):
20         print(f"Method: {self.method} - Iteration {i}")
21         sol, res = self.method(sol, res)
22         self.resnorm.append(NORM(res))

```

Code Listing 5: Solve method implementation

Although it is not the purpose of this work, the solve method is implemented in a way that allows the user to choose not only between the CG and PCG. Line

2 checks if a valid method is set. Lines 5 to 17 initialize the required vectors depending on the method set. If the method is the Conjugate Gradient, the search direction p_k is set. If a preconditioner is set, the preconditioned vector z is set.

Line 19 initializes the residual norm vector. At line 20, a loop is performed for the number of iterations set in the constructor and the method performs the solution of the linear system. The residual norm is stored at each iteration.

3.2 Conjugate Gradient Method

The ConjugateGradient method is called by the Solve method and solves the linear system. Code 6 shows the implementation of this method.

```

1 def ConjugateGradient(self, sol:np.array, res:np.array)->tuple[float, list[
    float]]:
2     alpha_k = INNER(res.T, res) / INNER(self.A.Multiply(self.p_k.T), self.
    p_k) if not self.preconditioner else INNER(self.res_k.T, self.z) / INNER
    (self.A.Multiply(self.p_k.T), self.p_k)
3
4     sol += alpha_k * self.p_k
5     res = self.res_k - alpha_k * self.A.Multiply(self.p_k)
6
7     if not self.preconditioner:
8         beta_k = INNER(res.T, res) / INNER(self.res_k.T, self.res_k)
9         self.p_k = res + beta_k * self.p_k
10
11    else:
12        self.z, _ = self.preconditioner(ZEROS(self.A.size), res)
13        beta_k = INNER(res.T, self.z) / INNER(self.res_k.T, self.z_k)
14
15        self.z_k = self.z.copy()
16        self.p_k = self.z + beta_k * self.p_k
17
18    self.res_k = res.copy()
19    return sol, res

```

Code Listing 6: ConjugateGradient method implementation

Aiming to avoid rewriting the code for the PCG method, the ConjugateGradient method is implemented in a way that allows the user to set a preconditioner. If a preconditioner is set, the method calculates the alpha and beta coefficients using the preconditioned vectors.

Line 2 calculates the alpha coefficient, considering whether a preconditioner is set. The solution is then updated in line 4 and the residual in line 5. If no preconditioner is set, the beta coefficient is calculated in line 8 and the direction is updated in line 10.

Conversely, if a preconditioner is set, the preconditioned vector is updated in line 13, the beta coefficient is calculated in line 14 and the direction is updated in line 18. The residual norm is evaluated in line 20 regardless of the preconditioner.

3.3 The Jacobi Preconditioner

Two preconditioners are implemented in this work: the Jacobi and the SSOR. The Jacobi preconditioner is implemented in Code 7.

```
1 def Jacobi(self, sol:np.array, res:np.array)->None:
2     M = self.A.GetDiagonal()
3     reslocal = res.copy()
4
5     dx = self.omega * np.divide(res, M)
6
7     reslocal -= self.A.Multiply(dx)
8
9     return sol + dx, reslocal
```

Code Listing 7: Jacobi preconditioner implementation

This preconditioner follows the idea behind the Jacobi method, evaluating the matrix M as the diagonal of the matrix A (line 2). The Solution is then updated by the division of the current residual and the matrix M (line 5). The

residual for the next step is obtained by subtracting the matrix-vector product of matrix A and the solution update (line 7).

3.4 The SSOR Preconditioner

The second preconditioner implemented is the SSOR. The SSOR's procedure employs the Gauss-Seidel method, performing both forward and backward iterations at each step. Code 8 shows the implementation of this method.

```

1 def GaussSeidelF(self, sol:np.array, res:np.array)->np.array:
2     dx = ZEROS(self.A.size)
3     reslocal = res.copy()
4
5     dx[0] = self.omega * reslocal[0] / self.A.FindAij(0, 0)
6
7     for i in range(1, self.A.size):
8         reslocal[i] -= self.A.InnerProductLowerRows(dx, i)
9         dx[i] += self.omega * reslocal[i] / self.A.FindAij(i, i)
10
11     reslocal = res - self.A.Multiply(dx)
12
13     return sol + dx, reslocal
14
15 def GaussSeidelB(self, sol:np.array, res:np.array)->np.array:
16     dx = ZEROS(self.A.size)
17     reslocal = res.copy()
18
19     dx[self.A.size-1] = self.omega * reslocal[self.A.size-1] / self.A.
20     FindAij(self.A.size-1, self.A.size-1)
21
22     for i in range(self.A.size-2, -1, -1):
23         reslocal[i] -= self.A.InnerProductUpperRows(dx, i)
24         dx[i] += self.omega * reslocal[i] / self.A.FindAij(i, i)
25
26     reslocal = res - self.A.Multiply(dx)
27
28     return sol + dx, reslocal

```

```
29 def SSOR(self, sol:np.array, res:np.array)->np.array:
30     sol, res = self.GaussSeidelF(sol, res)
31     sol, res = self.GaussSeidelB(sol, res)
32
33     return sol, res
```

Code Listing 8: SSOR preconditioner implementation

As shown in Code 8, the SSOR method calls the GaussSeidelF and GaussSeidelB, performing two iterations at each step (lines 29 to 33). The GaussSeidelF method does a forward iteration, updating and the residual equation-wise (lines 8 and 9). The GaussSeidelB method, on the other hand, performs a backward iteration (lines 18 and 19), updating the residual equation-wise from the last to the first equation.

The updating of the residual is done by the vector inner product with the lower and upper triangles of the matrix A. The solution is updated by the division of the residual and the diagonal of the matrix A. Both methods return the updated solution and residual. However, while the GaussSeidelF uses the initial solution and residual, the GaussSeidelB uses the updated solution and residual from the GaussSeidelF method (lines 30 and 31).

The next section presents the results obtained by the linear system given in class.

4 Comparison of the Conjugate Gradient Method

The main code used to solve the linear system using the CG and PCG methods is shown in Code 9.

```
1 from List3.SparseMatrix import SparseMatrix
2 from IndirectSolver import IndirectSolver
3
4 def ParseVector(file: str)->np.array:
```

```
5     with open(file, "r") as f:
6         vector = np.array([float(x) for x in f.read().split(",")])
7
8     return vector
9
10 def Main()->None:
11     matrix_file = "List4/matrix.dat"
12     rhs_file = "List4/rhs.dat"
13     results_file = "List4/Results.py"
14
15     niter = 500
16
17     matrix = SparseMatrix()
18     matrix.ParseFromFile(matrix_file)
19
20     rhs = ParseVector(rhs_file)
21
22     solver = IndirectSolver(matrix, rhs, niter = niter)
23
24     solver.Set_method("ConjugateGradient")
25     solver.Solve()
26     resnormCG = solver.resnorm
27
28     solver.Set_preconditioner("Jacobi")
29     solver.Solve()
30     resnormCGJ = solver.resnorm
31
32     solver.Set_preconditioner("SSOR")
33     solver.Solve()
34     resnormCGSSOR = solver.resnorm
35
36     with open(results_file, 'w') as f:
37         print(f"resnormCG = {str(resnormCG)}", file=f)
38         print(f"resnormCGJ = {str(resnormCGJ)}", file=f)
39         print(f"resnormCGSSOR = {str(resnormCGSSOR)}", file=f)
40
41 if __name__ == "__main__":
```

42

Main()

Code Listing 9: Main code to solve the linear system using the CG and PCG methods.

Lines 1 and 2 import the required packages (SparseMatrix, from list 3, and IndirectSolver). Lines 4 to 8 define a function to parse the rhs.dat file and assign these values to a vector.

Line 10 defines the main function. From line 11 to line 22, the sparse matrix object and the solver object are created. Line 24 defines the Conjugate Gradient as the method to solve the linear system. Three different analyses are performed: first without any preconditioner, second with the Jacobi preconditioner and third with the SSOR preconditioner (lines 25 to 34).

Due to the time to iteratively solve the linear system, the results are saved in a file called Results.py, which is later imported in the PlottingResults.py file. Figure 4.1 depicts the results of the convergence, using 500 iterations, for the three different preconditioners.




Figure 4.1: Convergence of the Conjugate Gradient method with different preconditioners.

As it's possible to see, the CG without a preconditioner did not converge to the solution. The Jacobi preconditioner improved the convergence and reached an error of 10^{-3} in 500 iterations. The SSOR preconditioner, however, was the most efficient, solving the system in 500 iterations with machine precision.

5 Conclusions

References

A GitHub Repository

The source code for this report and every code inhere mentioned can be found in the following GitHub repository: [CarlosPuga14/MetodosNumericos_-2024S1](#).